
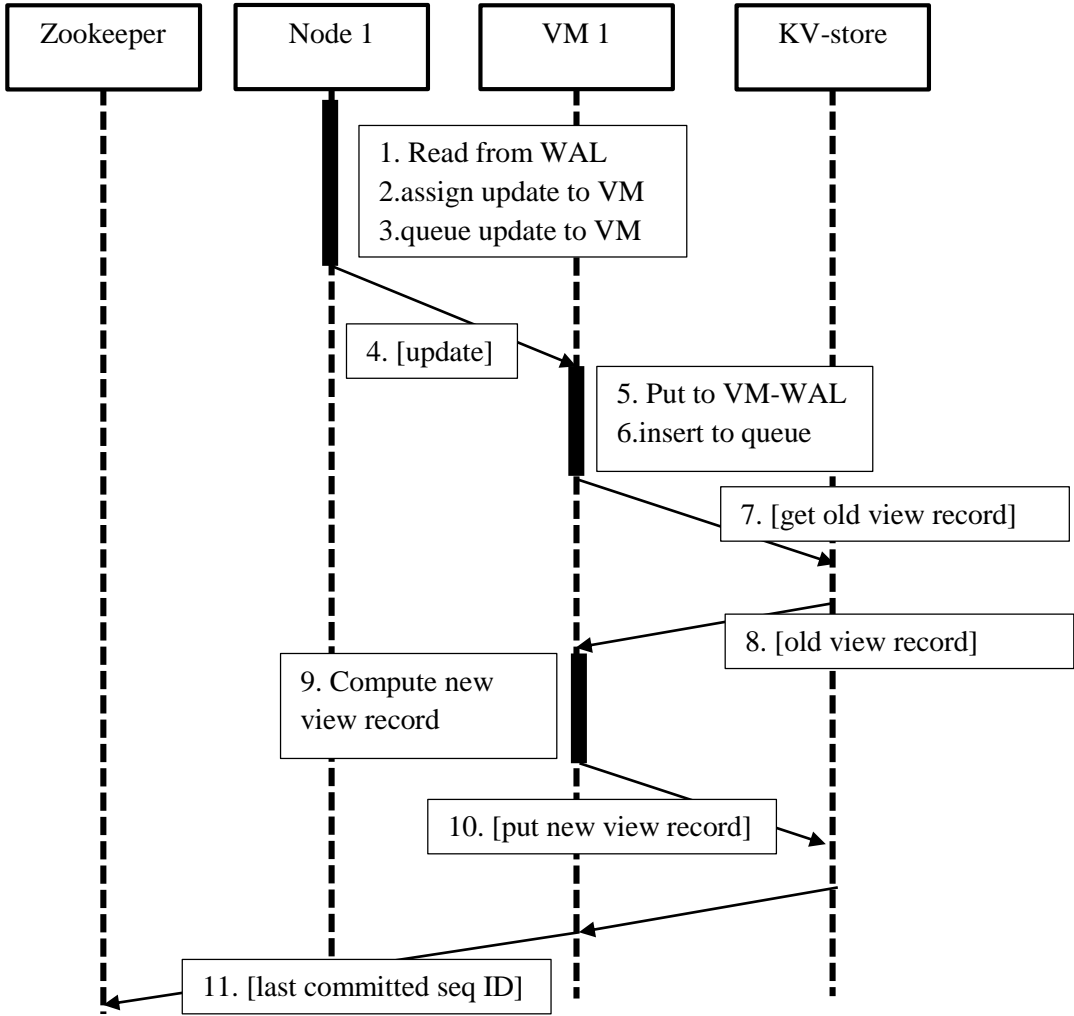


# Global Hashring

## Process update



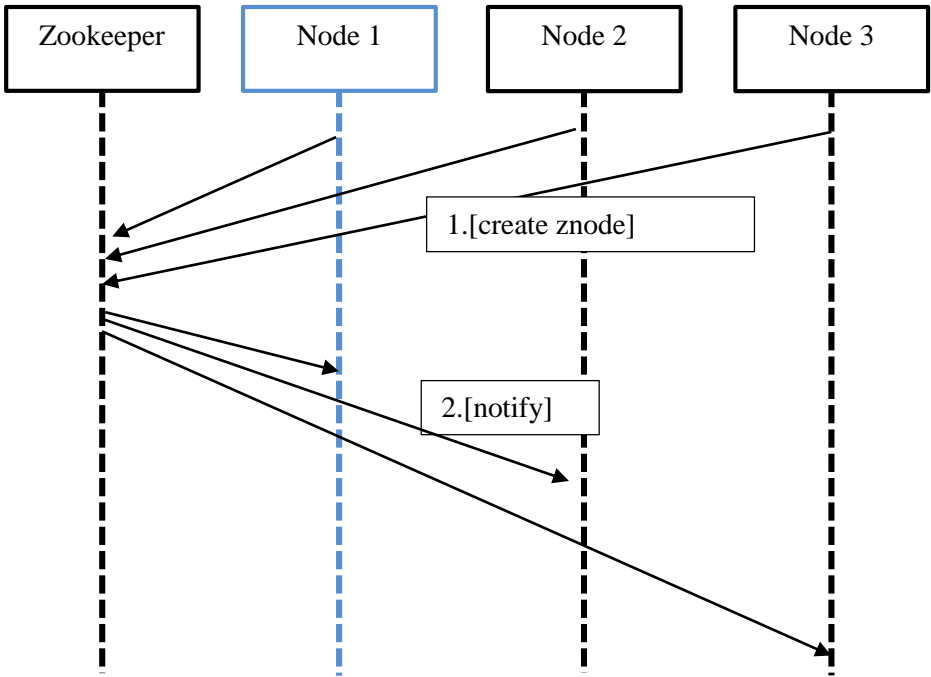

**Description:**  
Updates are sent through a different communication channel than messages. The update stream is continuously flowing from the database nodes to the view managers. One update passes the following stages.



**Algorithm:**  
All Nodes:[create znode] → Zookeeper  
Zookeeper:[election finished] → all nodes

# Global Hashring


## Elect Coordinator




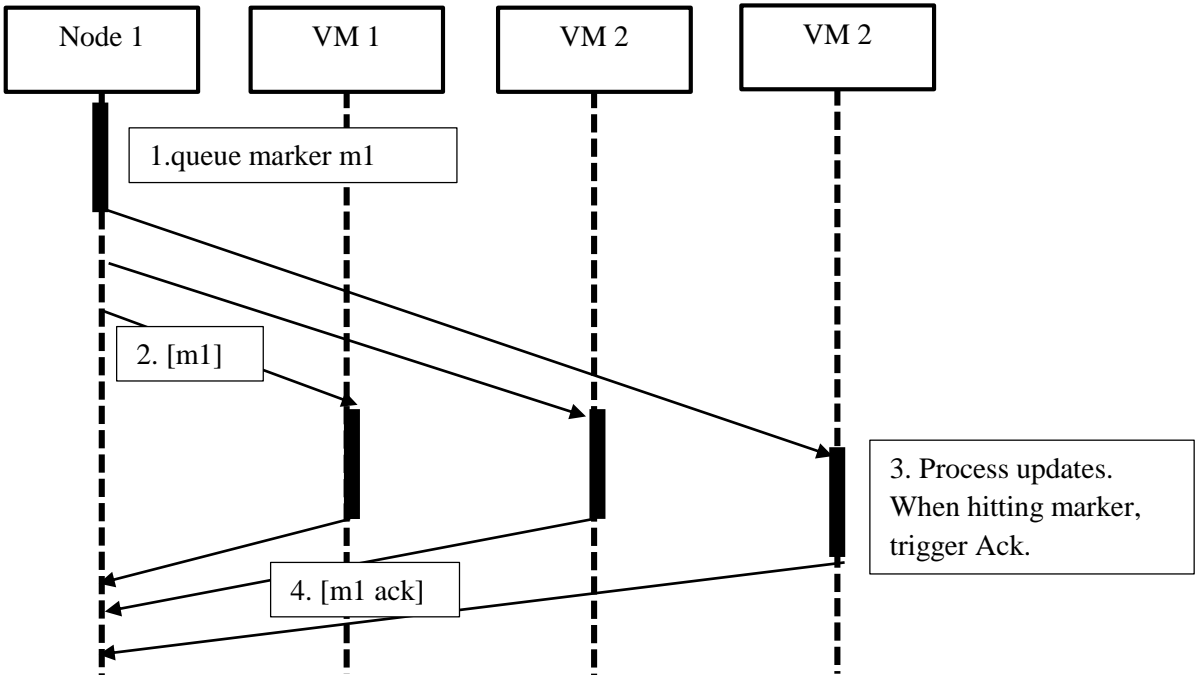
**Description:**  
The global hash ring architecture style of the VMS does not rely on a centralized coordinator. Nevertheless, the VMS needs to perform actions to load balance or recover the system. Without a coordinator these tasks can be assigned to one of the nodes. Using a leader election algorithm the node can be determined.

- Algorithm:**
1. All Nodes:[create znode] → Zookeeper
  2. Zookeeper:[election finished] → all nodes

## Send marker




**Description:**  
To perform actions like the assign and withdraw view managers in a dynamic environment, the Node requires information about whether a view manager has already processed a set of updates or not. For that reason markers are placed in the view manager’s queue like a normal update. From here they are sent to the view manager. At some point the VM hits the marker in the stream of updates. Then it




- Algorithm:**
- All Nodes:[create znode] → Zookeeper
  - Zookeeper:[election finished] → all nodes

# Global Hashring

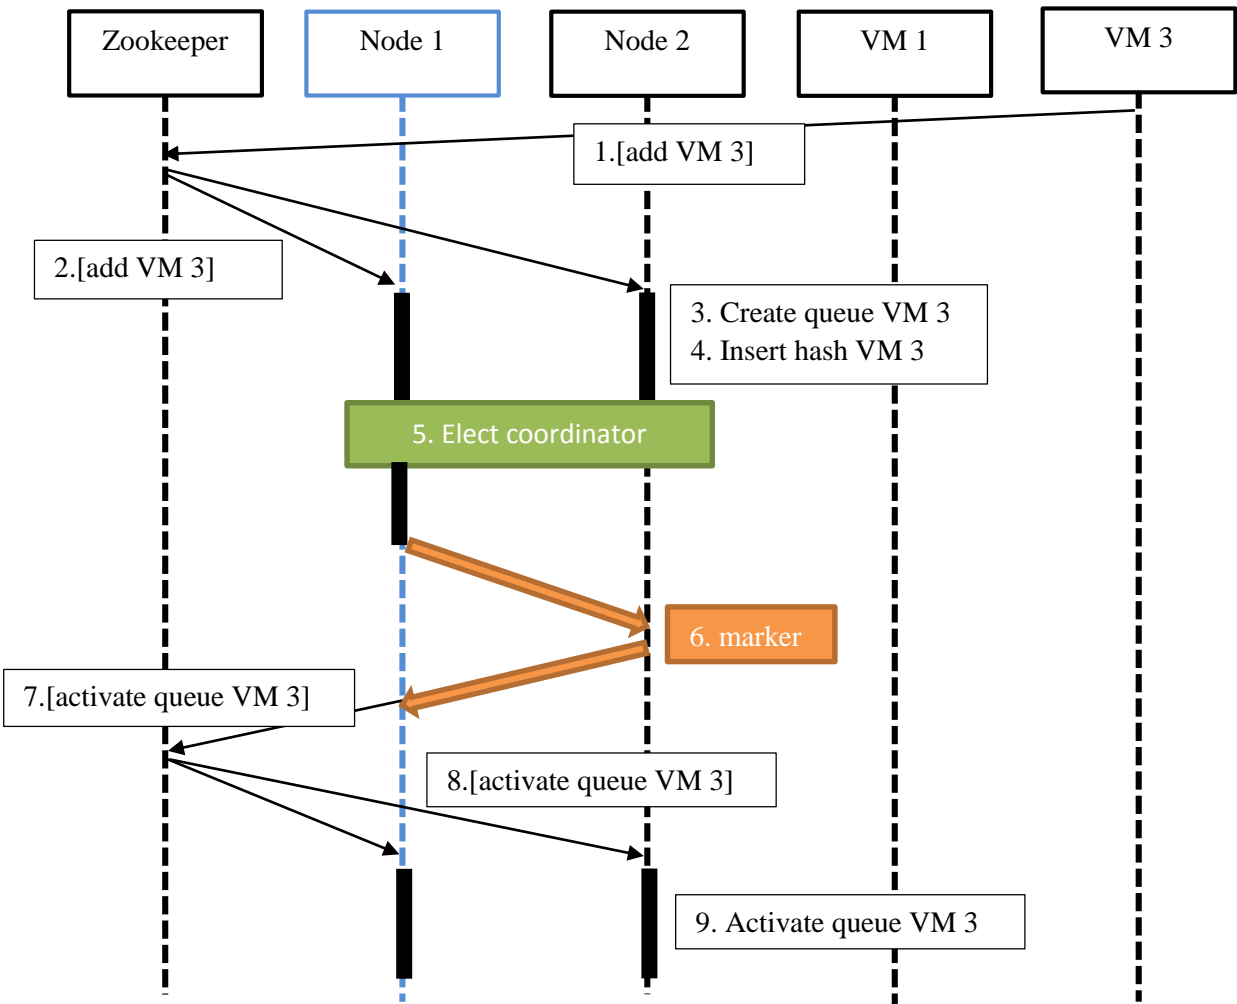
## Add View Manager



A View Manager VM 3 is added to VMS. Because the implementation uses a global hash ring, the VM needs to be added to the hash rings of all nodes. To preserve the time line, the queues of all VMs that lose a key-range are deactivated. **Alternative 1:** The nodes synchronize on Zookeeper, only one node receives task of activation. The Node sends a marker to the new VM. After receiving the Ack, the node reactivates the stopped VMs and activates the new VM. **Alternative 2:** All nodes send a marker to the new VM. On acknowledgement, the stopped VMs are reactivated and the new VM is activated.

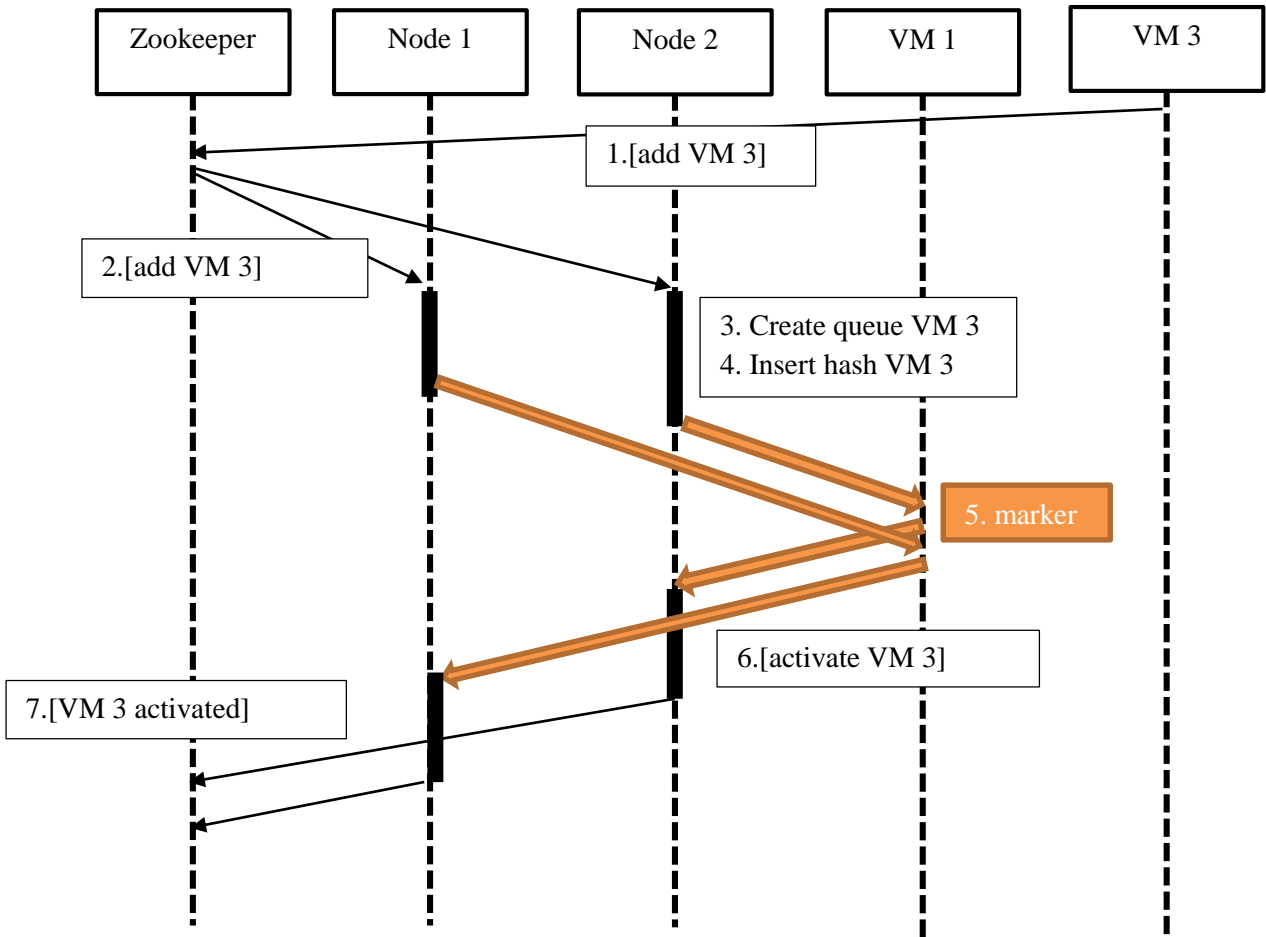


### Alternative 1 (single node activate):




1. New VM: Send [add new VM] → Zookeeper
2. Zookeeper: [new VM added] → all nodes
3. All nodes: Create queue
4. All nodes: Insert hash of new VM
5. Coordinator: Send [marker] to all VMs that lose key range
6. VM: Send [ack] to coordinator
7. Coordinator: Send [activate new VM] to all nodes

### Alternative 2 (all nodes activate):

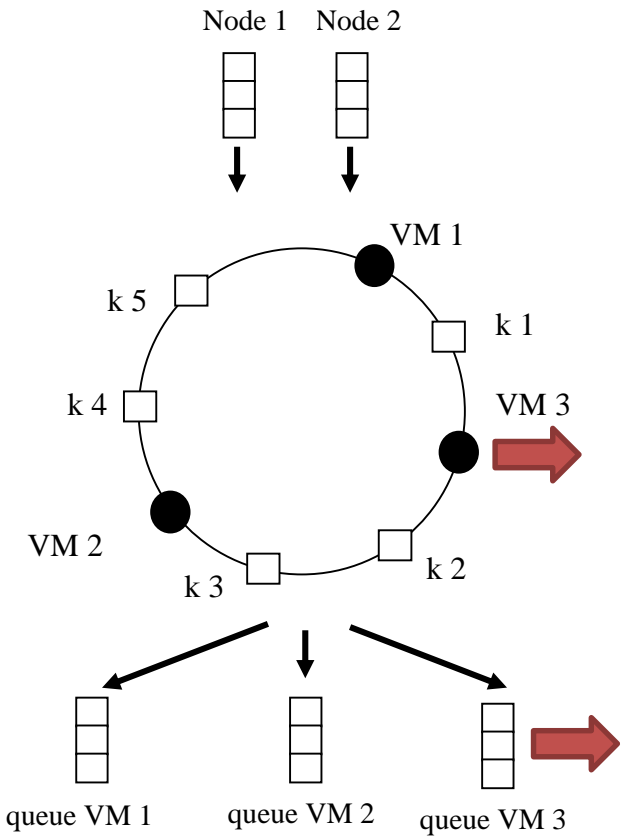


# Global Hashring

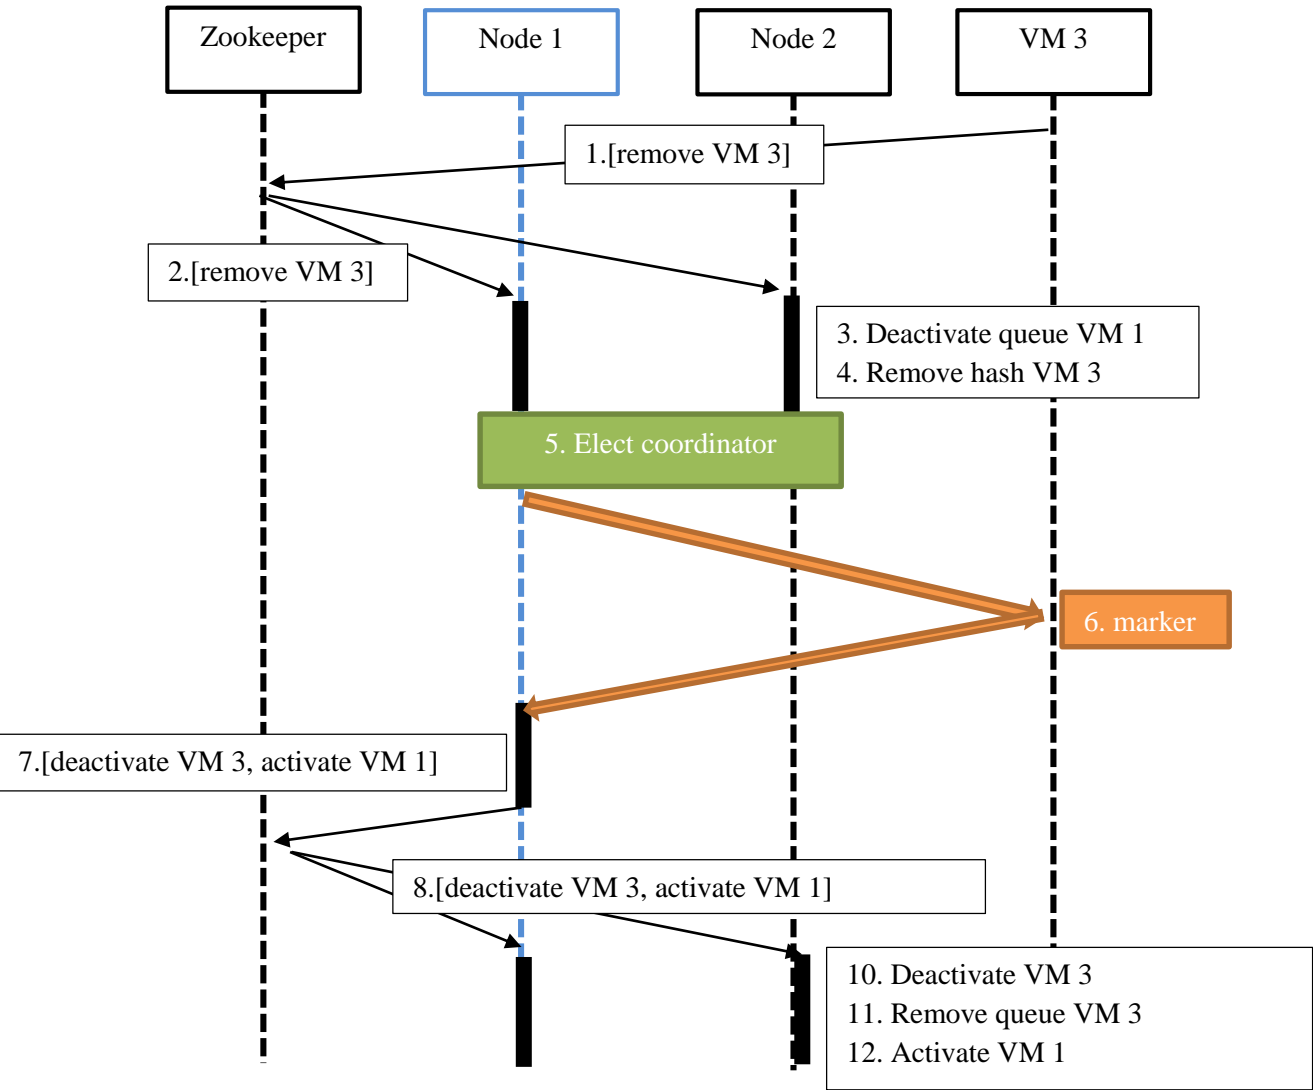
## Remove View Manager



- View Manager is added to VMS
- Needs to be removed globally (from hash rings of all the nodes)
- Timeline consistency has to be preserved
- Marker is sent to the view manager that is removed

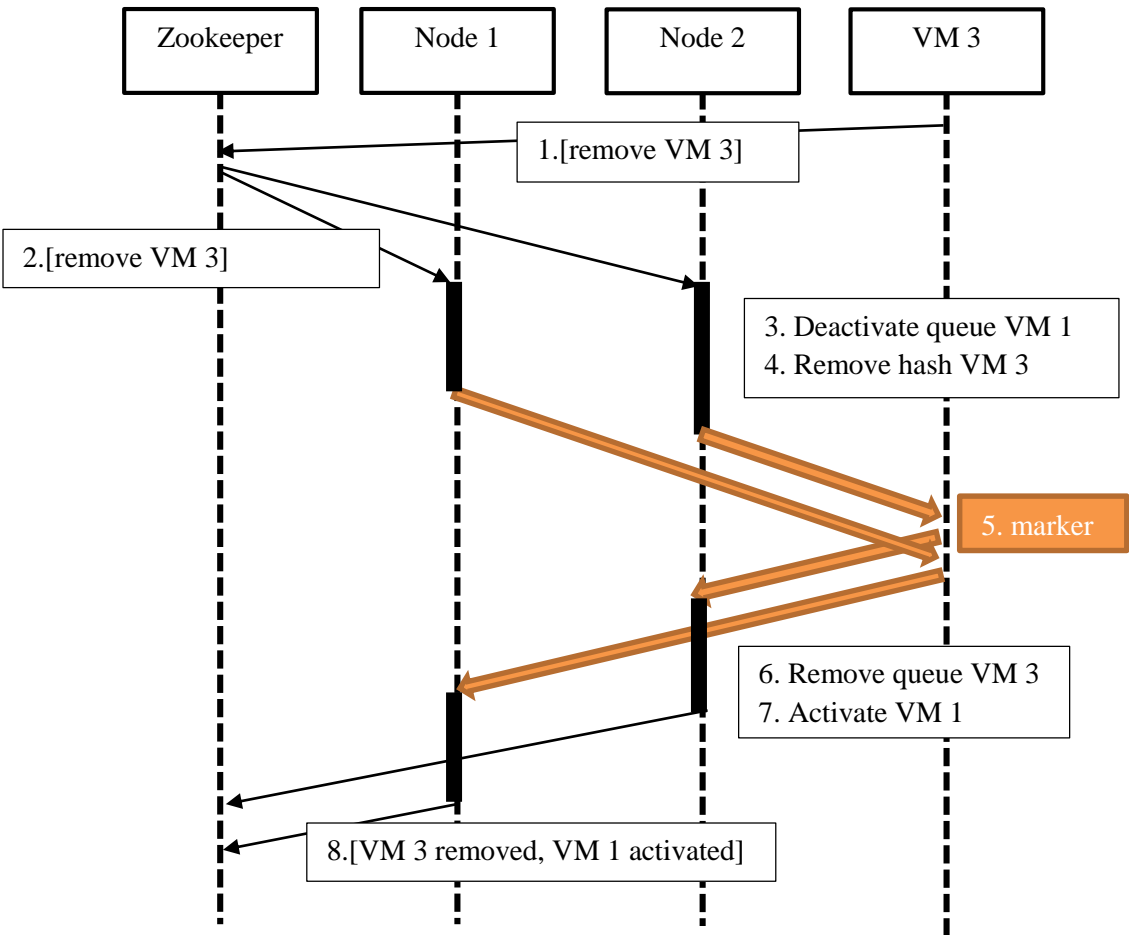


### Alternative 1 (with coordinator):




1. Coordinator: Send deactivate VM 1 + remove VM 3 from hashring to all nodes
2. All nodes: Deactivate queue VM 1
3. All nodes: Remove VM 3 from hash ring
4. All nodes: send [ack]
5. Coordinator: Queue marker to VM 3
6. Coordinator: Wait for VM 3 to ack
7. Coordinator: Send activate VM 1 to all nodes

### Alternative 2 (without coordinator):

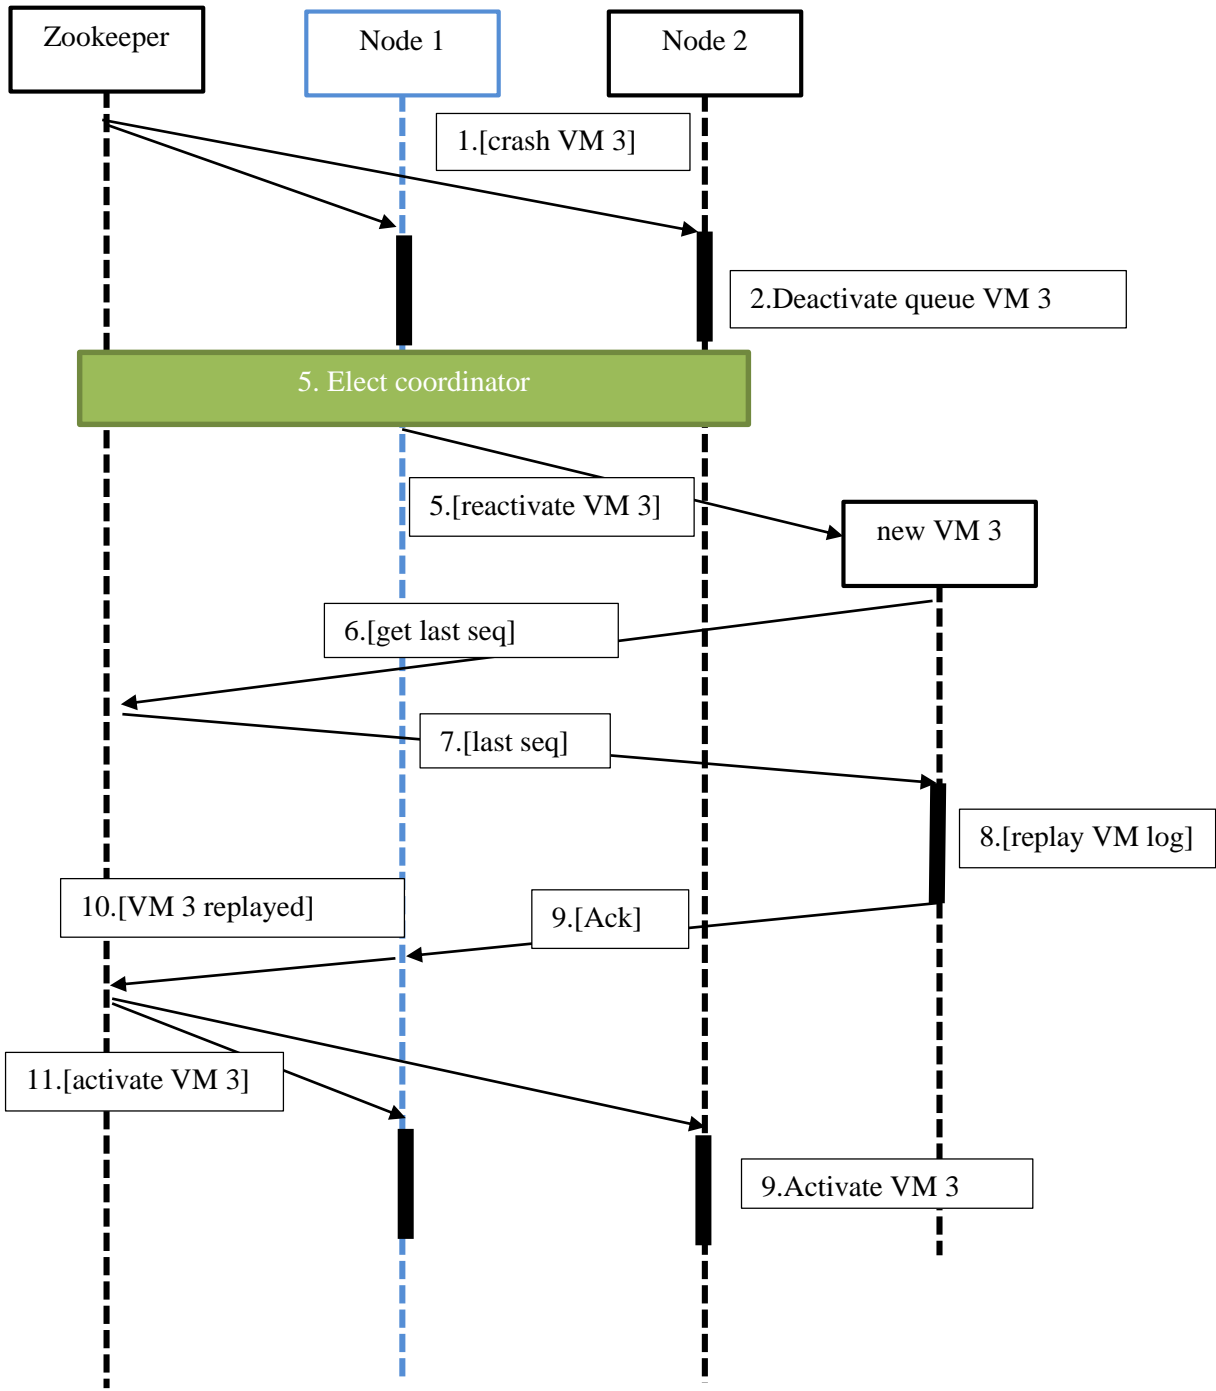



# Global Hashring

## Crash View Manager

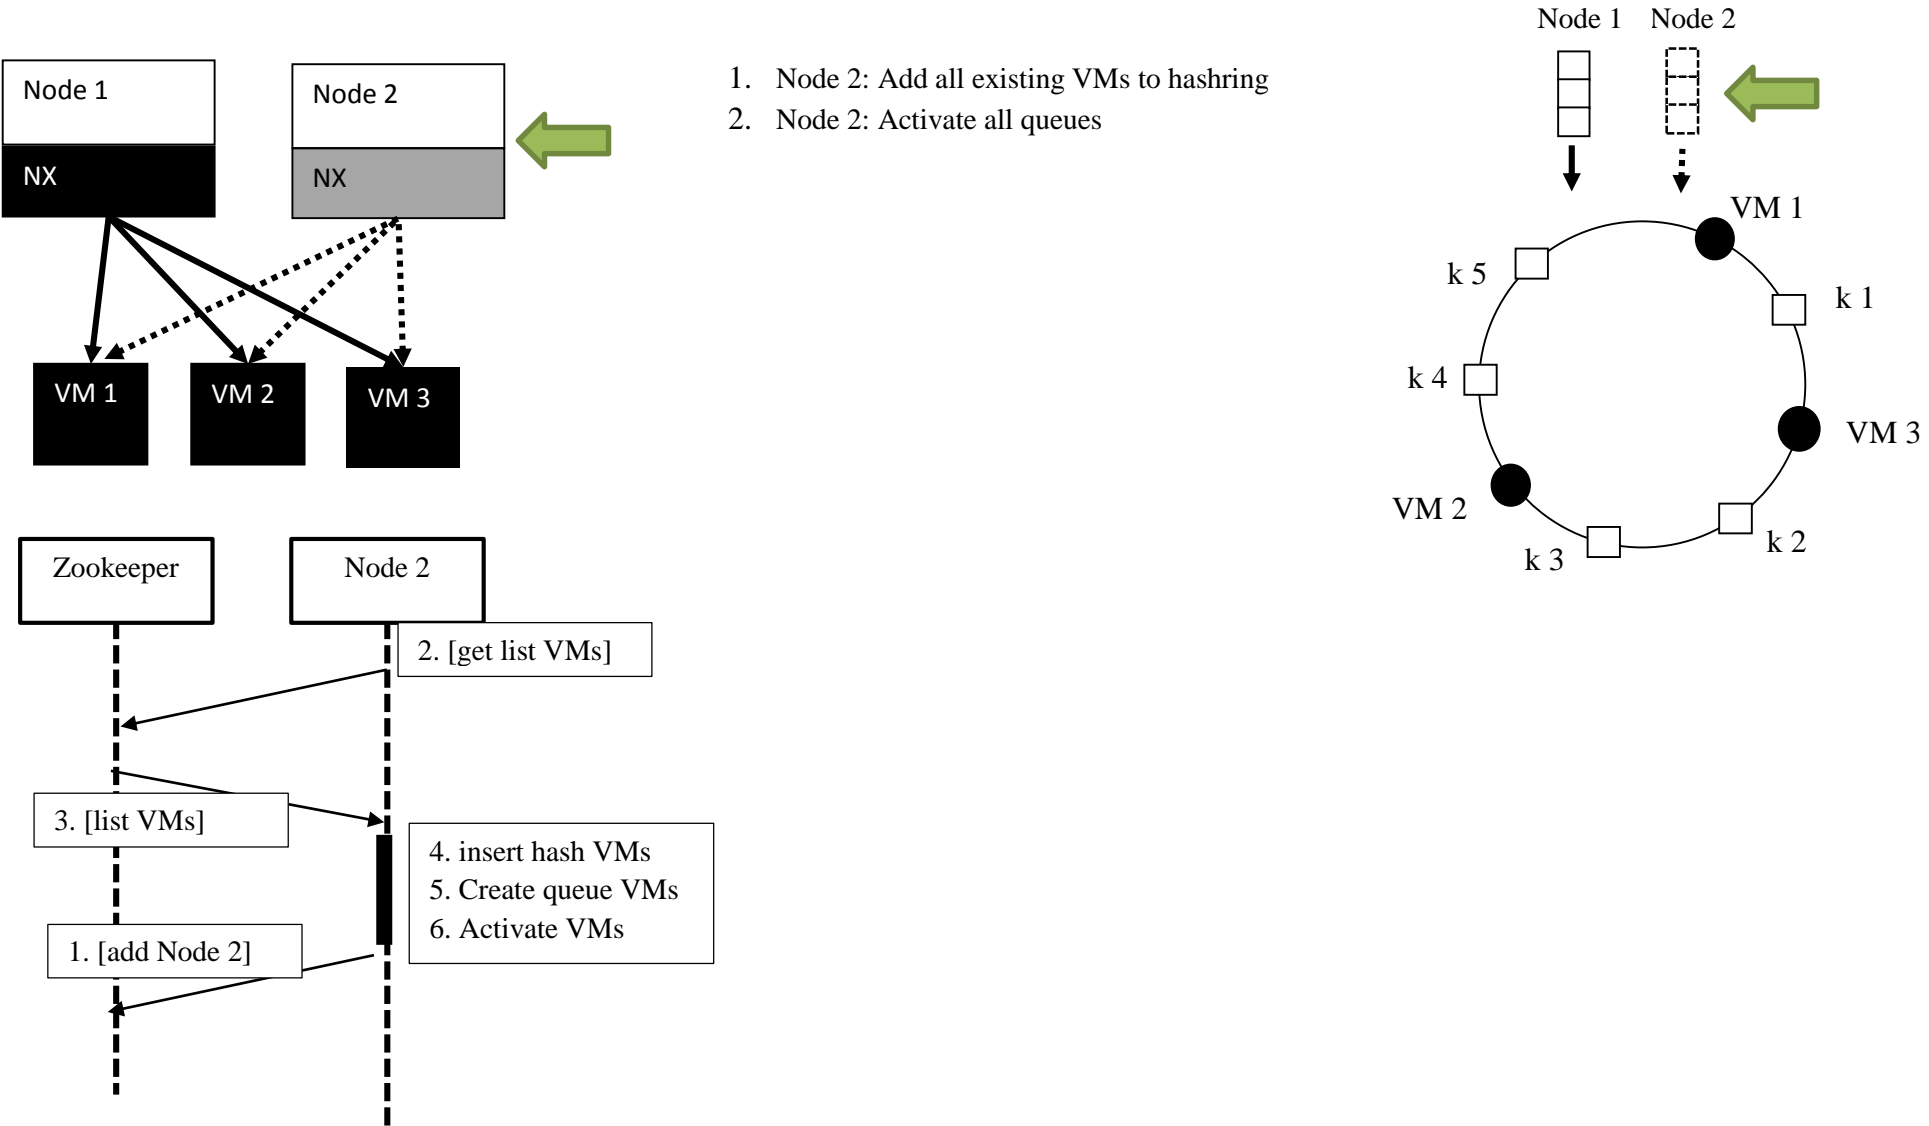


**Description:**  
In case a view manager crashes a node is elected that performs the recovery steps. The Node restarts the crashed view manager (on a different machine). On startup, the view manager. In the first case the View Manager needs only to keep track of its *last committed* sequence ID. The WAL is replayed from this point

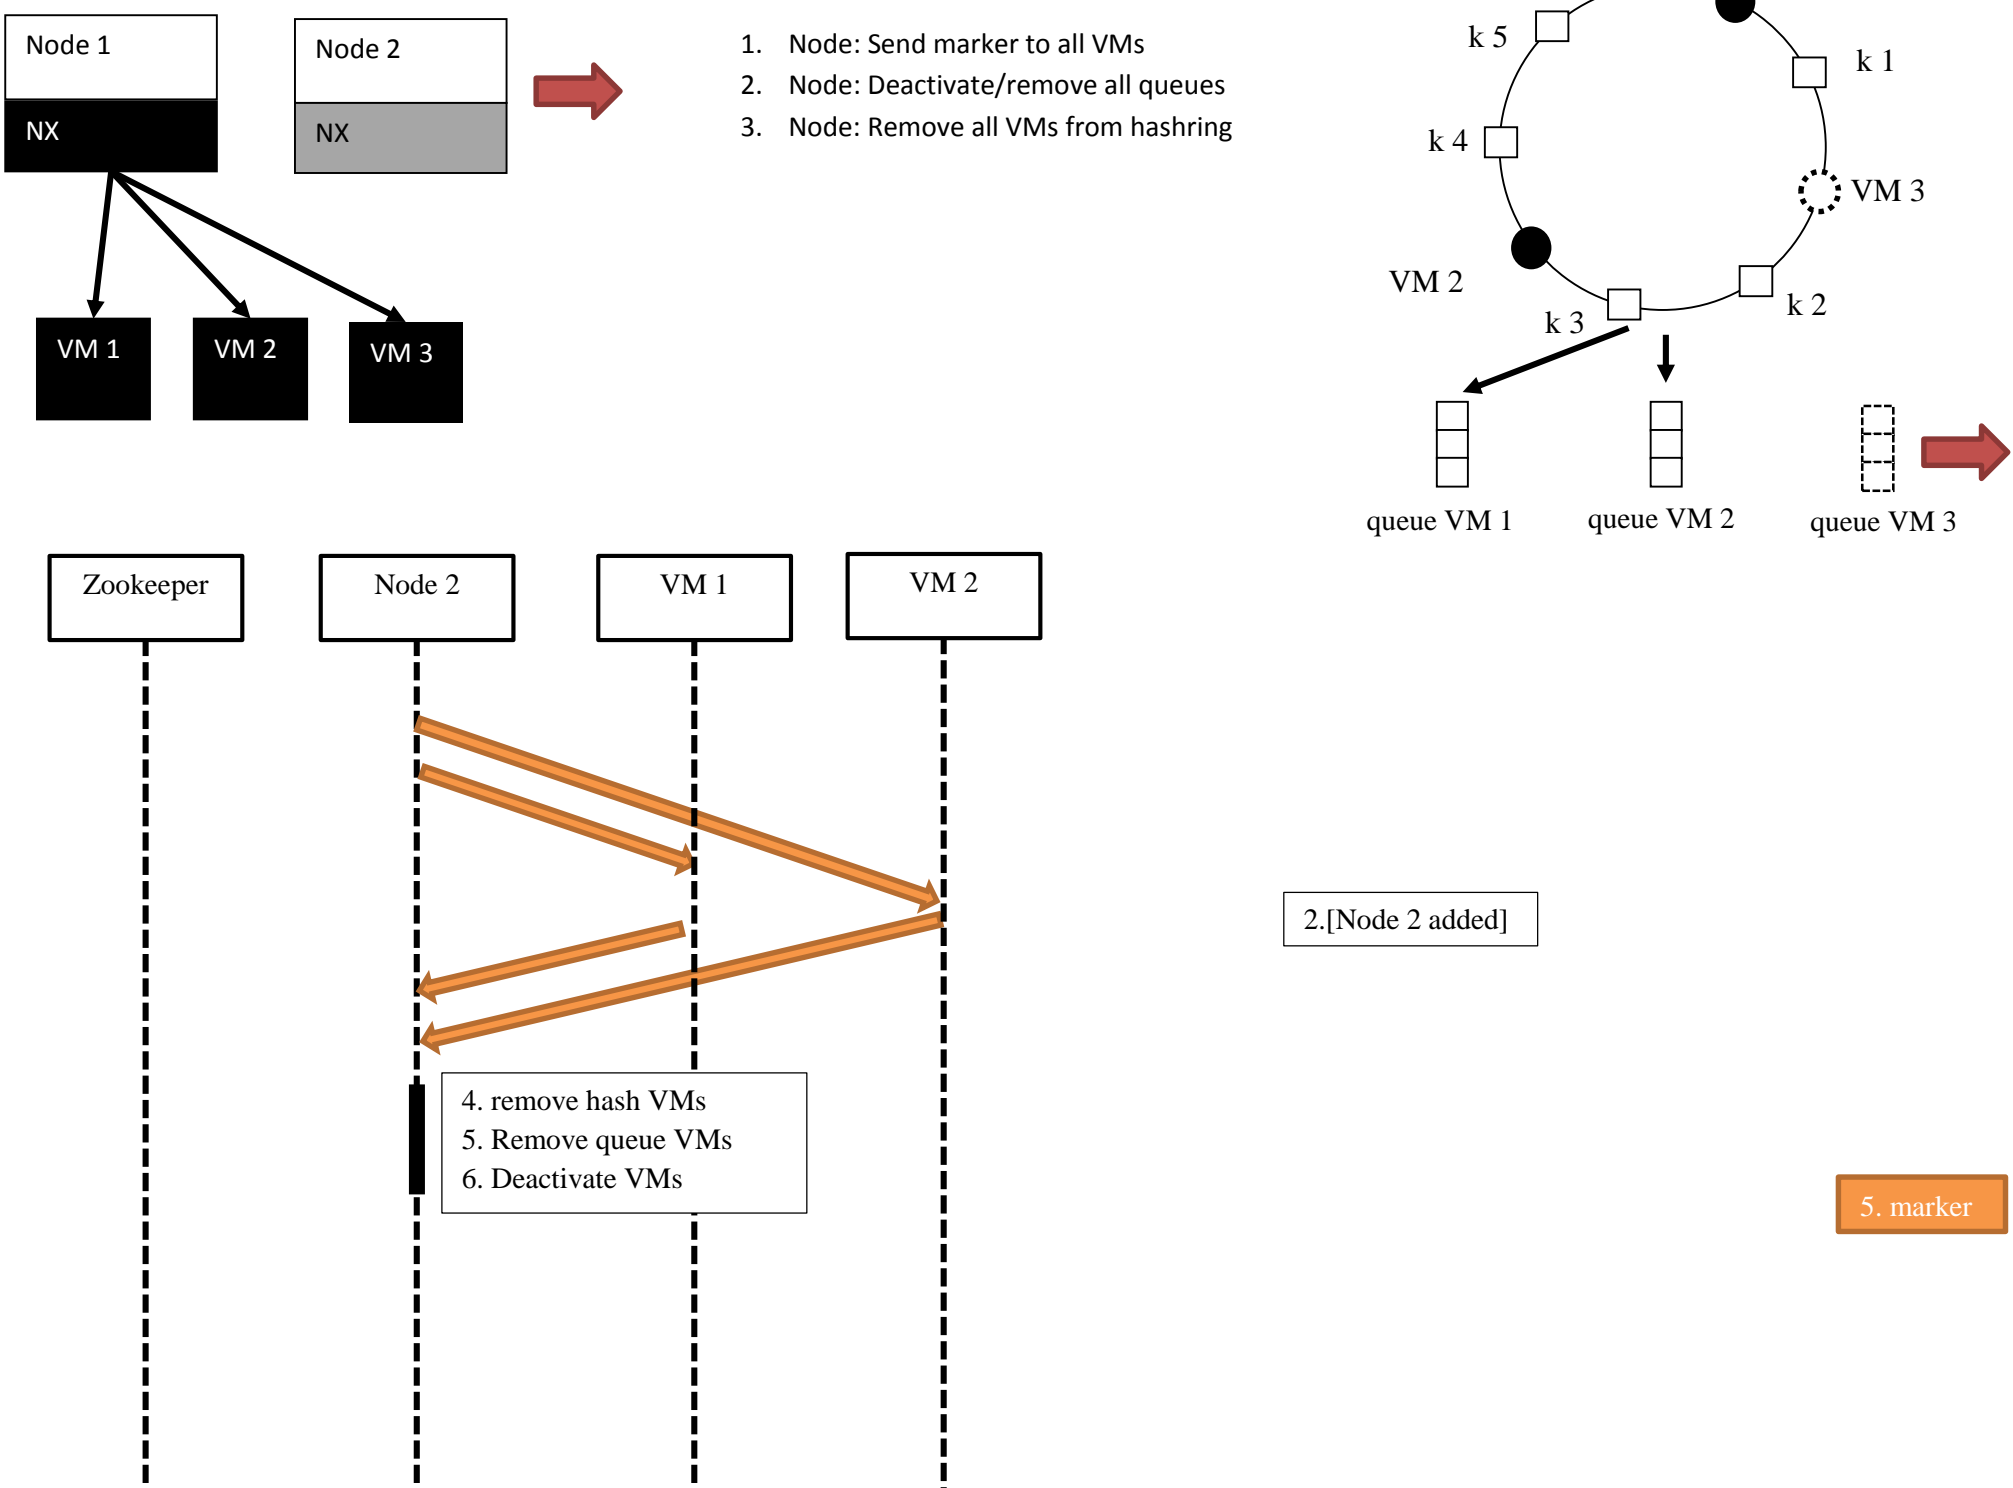


1. Coordinator: Send [deactivate queue of crashed VM] to all nodes
2. All nodes: Deactivate queue of crashed VM
3. All nodes: Send [ack]
4. Coordinator: Start new Vm
5. Coordinator: Send [replay log of crashed VM] to new VM
6. New VM: replay log of crashed VM
7. New VM: send [ack]
8. Coordinator: Send [substitute ip address] to new VM

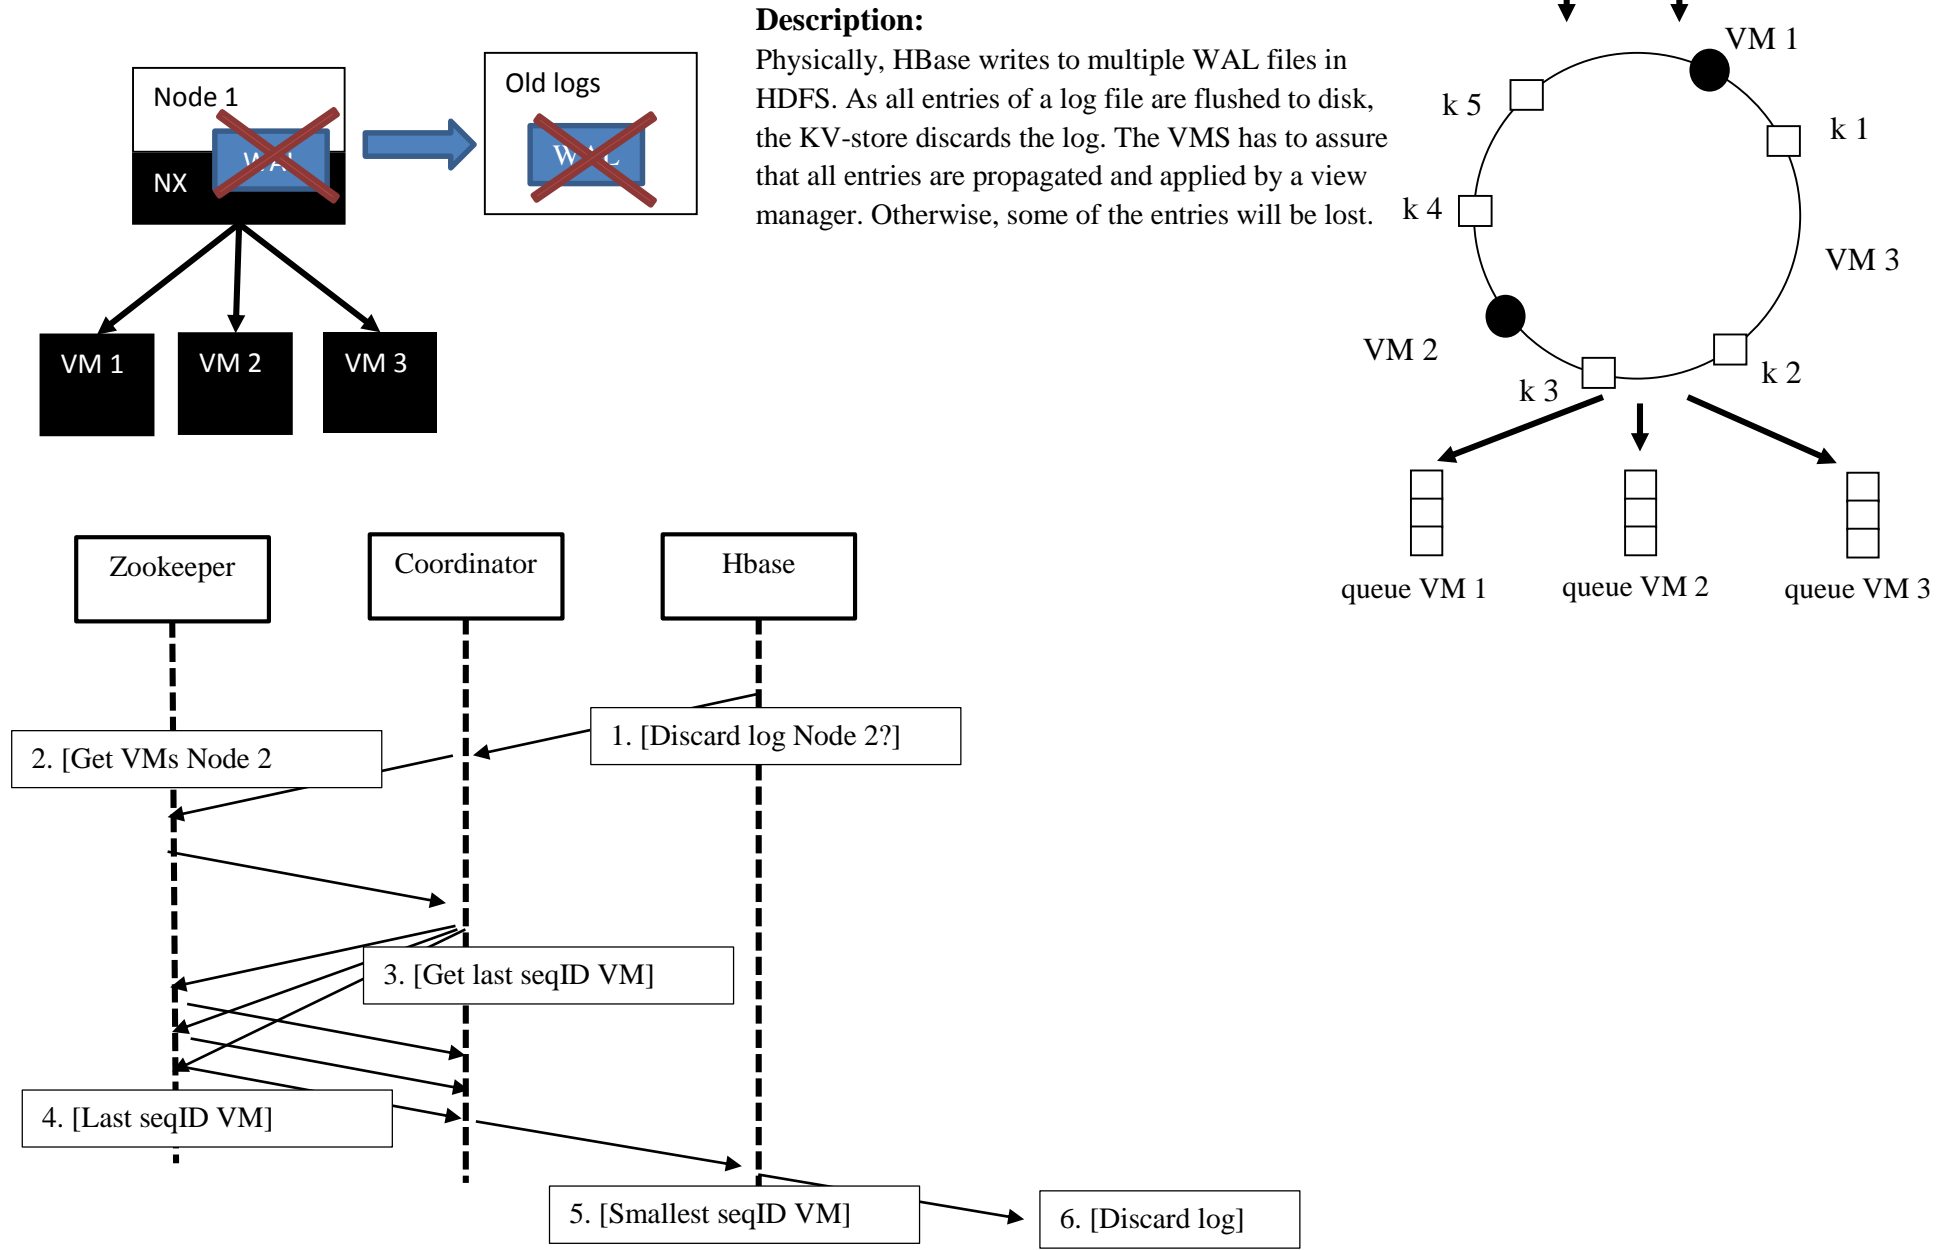
Add Node



Remove Node



Discard logs



Description:

Physically, HBase writes to multiple WAL files in HDFS. As all entries of a log file are flushed to disk, the KV-store discards the log. The VMS has to assure that all entries are propagated and applied by a view manager. Otherwise, some of the entries will be lost.

Zookeeper

Coordinator

Hbase

2. [Get VMs Node 2]

1. [Discard log Node 2?]

3. [Get last seqID VM]

4. [Last seqID VM]

5. [Smallest seqID VM]

6. [Discard log]

Node 1

Node 2

VM 1

VM 2

VM 3

k 5

k 4

k 3

k 2

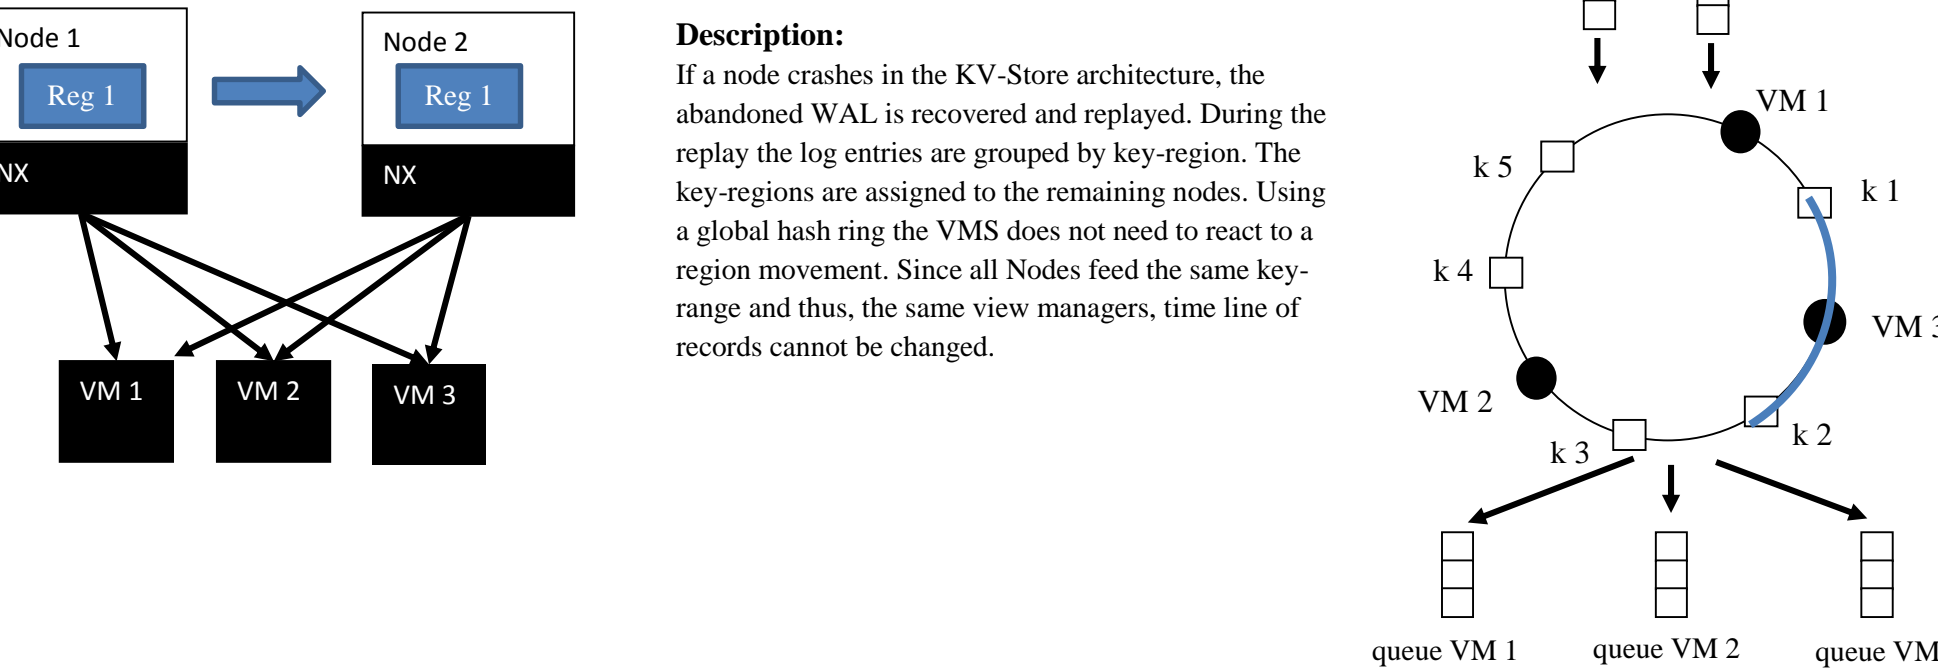
k 1

queue VM 1

queue VM 2

queue VM 3

Move Region



Crash Node (move region + discard logs )

